

Go (Golang) Distributed Rate Limiter - Token Bucket Algorithm

```
// Package ratelimit provides a concurrent-safe distributed rate limiter
```

```
// implementation using token bucket algorithm backed by Redis.
```

```
// It supports context cancellation, pluggable storage backends, and
```

```
// extensible limiter policies.
```

```
//
```

```
// This package can be used to throttle API calls, limit resource usage,
```

```
// or implement fair usage policies in distributed systems.
```

```
package ratelimit
```

```
import (
```

```
    "context"
```

```
    "errors"
```

```
    "fmt"
```

```
    "sync"
```

```
    "time"
```

```
    "github.com/go-redis/redis/v8"
```

```
)
```

```
// Limiter defines the interface for any rate limiter implementation.
```

```

// It supports Acquire and Release semantics for tokens.

type Limiter interface {
    // Acquire blocks until a token is available or context is canceled.
    Acquire(ctx context.Context, key string) error

    // TryAcquire tries to acquire a token immediately, returns false if unavailable.
    TryAcquire(ctx context.Context, key string) (bool, error)

    // Release returns a token back to the bucket.
    Release(ctx context.Context, key string) error
}

// ErrRateLimitExceeded indicates the rate limit was exceeded.
var ErrRateLimitExceeded = errors.New("rate limit exceeded")

// Storage defines the interface for backend storage supporting
// atomic operations required by the limiter.
type Storage interface {
    // Increment increases token count and returns the new count.
    Increment(ctx context.Context, key string, delta int64, expire time.Duration)
    (int64, error)

    // Decrement decreases token count and returns the new count.
    Decrement(ctx context.Context, key string, delta int64) (int64, error)
}

```

```

    // Get returns the current token count for a key.
    Get(ctx context.Context, key string) (int64, error)
}

// redisStorage implements Storage using Redis as backend.
type redisStorage struct {
    client *redis.Client
}

func newRedisStorage(client *redis.Client) *redisStorage {
    return &redisStorage{client: client}
}

func (r *redisStorage) Increment(ctx context.Context, key string, delta int64, expire
time.Duration) (int64, error) {
    pipe := r.client.TxPipeline()
    incr := pipe.IncrBy(ctx, key, delta)
    pipe.Expire(ctx, key, expire)
    _, err := pipe.Exec(ctx)
    if err != nil {
        return 0, err
    }
    return incr.Val(), nil
}

```

```

func (r *redisStorage) Decrement(ctx context.Context, key string, delta int64) (int64,
error) {

    newVal, err := r.client.DecrBy(ctx, key, delta).Result()

    if err != nil {

        return 0, err

    }

    if newVal < 0 {

        // Optionally reset to zero to avoid negative tokens

        _ = r.client.Set(ctx, key, 0, 0).Err()

        return 0, ErrRateLimitExceeded

    }

    return newVal, nil

}

```

```

func (r *redisStorage) Get(ctx context.Context, key string) (int64, error) {

    val, err := r.client.Get(ctx, key).Int64()

    if err == redis.Nil {

        return 0, nil // Key not found implies zero tokens

    }

    return val, err

}

```

// TokenBucketLimiter is a distributed token bucket rate limiter.

// Tokens are replenished at a fixed interval up to a max capacity.

```

type TokenBucketLimiter struct {

```

```

    storage  Storage

    capacity int64    // max tokens in bucket

    refillRate int64    // tokens added per refill interval

    refillIntv time.Duration // refill interval duration


    mu      sync.Mutex

    closeCh  chan struct{ }

    closeOnce sync.Once
}

// NewTokenBucketLimiter constructs a new limiter with specified params.

func NewTokenBucketLimiter(storage Storage, capacity, refillRate int64, refillIntv
time.Duration) *TokenBucketLimiter {

    limiter := &TokenBucketLimiter{

        storage:  storage,

        capacity:  capacity,

        refillRate: refillRate,

        refillIntv: refillIntv,

        closeCh:  make(chan struct{ }),

    }

    go limiter.refillWorker()

    return limiter
}

// refillWorker replenishes tokens for all keys periodically.

```

// In a real-world system, you'd track active keys more efficiently.

```
func (l *TokenBucketLimiter) refillWorker() {  
    ticker := time.NewTicker(l.refillIntv)  
    defer ticker.Stop()  
  
    for {  
        select {  
        case <-ticker.C:  
            // Ideally: refill all keys in storage (implementation depends on  
backend)  
            // Here omitted for brevity as Redis SCAN required  
        case <-l.closeCh:  
            return  
        }  
    }  
}
```

// Close stops the background refill worker.

```
func (l *TokenBucketLimiter) Close() {  
    l.closeOnce.Do(func() {  
        close(l.closeCh)  
    })  
}
```

// Acquire blocks until a token is available or context canceled.

```

func (l *TokenBucketLimiter) Acquire(ctx context.Context, key string) error {
    for {
        ok, err := l.TryAcquire(ctx, key)

        if err != nil {
            return err
        }

        if ok {
            return nil
        }

        select {
        case <-ctx.Done():
            return ctx.Err()

        case <-time.After(l.refillIntv):
            // Wait and retry
        }
    }
}

```

// TryAcquire tries to acquire a token immediately.

```

func (l *TokenBucketLimiter) TryAcquire(ctx context.Context, key string) (bool, error) {
    l.mu.Lock()

    defer l.mu.Unlock()

    // Decrement token count atomically

```

```

    count, err := l.storage.Decrement(ctx, key, 1)

    if err == ErrRateLimitExceeded {
        return false, nil
    }

    if err != nil {
        return false, err
    }

    if count < 0 {
        return false, nil
    }

    return true, nil
}

// Release returns a token to the bucket.

func (l *TokenBucketLimiter) Release(ctx context.Context, key string) error {
    l.mu.Lock()
    defer l.mu.Unlock()

    count, err := l.storage.Increment(ctx, key, 1, 0)

    if err != nil {
        return err
    }

    if count > l.capacity {
        // Optionally cap tokens at capacity

```



```

        // Reset to capacity value

        // This ensures bucket doesn't overflow

        // For demo, ignoring atomicity tradeoffs
    }

    return nil
}

// --- Example Usage ---

// Example creates a limiter and tries to acquire tokens concurrently.
func Example() {
    ctx := context.Background()

    redisClient := redis.NewClient(&redis.Options{
        Addr: "localhost:6379",
    })

    storage := newRedisStorage(redisClient)

    limiter := NewTokenBucketLimiter(storage, 10, 2, time.Second*1)
    defer limiter.Close()

    keys := []string{"api_user_1", "api_user_2"}

    var wg sync.WaitGroup
    for _, key := range keys {

```

```

    wg.Add(1)
    go func(k string) {
        defer wg.Done()
        for i := 0; i < 15; i++ {
            err := limiter.Acquire(ctx, k)
            if err != nil {
                fmt.Printf("Failed to acquire token for %s: %v\n",
k, err)
                continue
            }
            fmt.Printf("Acquired token for %s\n", k)
            time.Sleep(100 * time.Millisecond)
            err = limiter.Release(ctx, k)
            if err != nil {
                fmt.Printf("Failed to release token for %s: %v\n", k,
err)
            }
        }
    }(key)
}
wg.Wait()
}

```

// --- Testing ---

// Below is a simplified unit test example.

// In production, place tests in separate _test.go files.

```
func testTokenBucketLimiter() error {  
    ctx := context.Background()  
  
    // Mock storage for unit test - in-memory storage implementation  
    memStore := newInMemoryStorage()  
  
    limiter := NewTokenBucketLimiter(memStore, 5, 1, time.Second)  
    defer limiter.Close()  
  
    key := "test_user"  
  
    // Preload tokens  
    _, err := memStore.Increment(ctx, key, 5, 0)  
    if err != nil {  
        return err  
    }  
  
    // Acquire tokens one by one  
    for i := 0; i < 5; i++ {  
        ok, err := limiter.TryAcquire(ctx, key)  
        if err != nil {  
            return fmt.Errorf("try acquire failed: %w", err)  
        }  
    }  
}
```

```

        if !ok {
            return fmt.Errorf("expected token available at iteration %d", i)
        }
    }

    // Next acquire should fail
    ok, err := limiter.TryAcquire(ctx, key)

    if err != nil {
        return err
    }

    if ok {
        return errors.New("expected rate limit exceeded")
    }

    return nil
}

// --- In-memory storage implementation for testing ---

type inMemoryStorage struct {
    mu    sync.Mutex
    store map[string]int64
}

func newInMemoryStorage() *inMemoryStorage {

```

```

    return &inMemoryStorage{
        store: make(map[string]int64),
    }
}

```

```

func (m *inMemoryStorage) Increment(ctx context.Context, key string, delta int64,
expire time.Duration) (int64, error) {
    m.mu.Lock()
    defer m.mu.Unlock()
    m.store[key] += delta
    return m.store[key], nil
}

```

```

func (m *inMemoryStorage) Decrement(ctx context.Context, key string, delta int64)
(int64, error) {
    m.mu.Lock()
    defer m.mu.Unlock()
    val, ok := m.store[key]
    if !ok {
        return 0, ErrRateLimitExceeded
    }
    val -= delta
    if val < 0 {
        m.store[key] = 0
        return 0, ErrRateLimitExceeded
    }
}

```

```
    m.store[key] = val  
    return val, nil  
}
```

```
func (m *inMemoryStorage) Get(ctx context.Context, key string) (int64, error) {  
    m.mu.Lock()  
    defer m.mu.Unlock()  
    val, ok := m.store[key]  
    if !ok {  
        return 0, nil  
    }  
    return val, nil  
}
```